

In [2]:

```
import pandas as pd
import numpy as np

from datetime import datetime
from sklearn.externals import joblib

import itertools
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.linear_model import LinearRegression, LogisticRegression, RidgeCV
from sklearn.model_selection import train_test_split, learning_curve, cross_val_score
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB, MultinomialNB

from sklearn import svm
from sklearn.svm import SVC

from sklearn.cross_validation import KFold
from sklearn.dummy import DummyClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, RandomForestRegressor, GradientBoostingClassifier, GradientBoostingRegressor

from sklearn.preprocessing import label_binarize, LabelEncoder
from sklearn.preprocessing import PolynomialFeatures

from sklearn.pipeline import make_pipeline

from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score, roc_auc_score
from sklearn.metrics import confusion_matrix, precision_recall_fscore_support, classification_report

%matplotlib inline
pd.set_option('display.max_columns', 400)
```

```
/Users/Sonal/anaconda/envs/py35/lib/python3.5/site-packages/sklearn/cross_validation.py:44: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model_selection module into which all the refactored classes and functions are moved. Also note that the interface of the new CV iterators are different from that of this module. This module will be removed in 0.20.
```

```
"This module will be removed in 0.20.", DeprecationWarning)
```

In [2]:

```
ic_df = pd.read_csv("ic_3mill.csv")
ic_df.head()
```

...

In []:

In [3]:

```
#changing the order of the columns  
ic_df = ic_df[['order_id','user_id','department_id','department','aisle_id','aisle'],  
ic_df.head()
```

```
...
```

In [4]:

```
#removed duplicate columns  
ic_df.head()
```

```
...
```

In [5]:

```
len(ic_df['department_id'].unique())
```

```
...
```

In [6]:

```
ic_df.shape
```

```
...
```

In [7]:

```
#checking which column have null values  
ic_df.isnull().any()
```

```
...
```

In [8]:

```
ic_df['days_since_prior'].isnull().sum().sum()
```

```
...
```

In [9]:

```
ic_df[pd.isnull(ic_df['days_since_prior'])]
```

```
...
```

In [10]:

```
ic_df['days_since_prior'].value_counts()
```

```
...
```

Orders per user

In [11]:

```
#orders per user  
ic_df.groupby('user_id',as_index=False)['order_number'].max()
```

...

In [12]:

```
ic_df.groupby('user_id',as_index=False)['order_number'].max().describe()
```

...

In [20]:

```
#Understanding the distributions - mean, min, max, median for both the variables to  
#orders_per_user_df.order_id.describe()
```

In [14]:

```
#Avg Items in Cart per Order Per User  
avg_items_per_order_per_user_df = ic_df.groupby(['user_id','order_id'],as_index=False)  
avg_items_per_order_per_user_df
```

...

items_per_order_per_user

In [15]:

```
#6.36 items per order per user  
items_per_order_per_user_df = avg_items_per_order_per_user_df.groupby(['user_id'],as_index=False)  
items_per_order_per_user_df.loc[:, 'add_to_cart_order'] = np.ceil(items_per_order_per_user_df['items_per_order_per_user'])  
items_per_order_per_user_df
```

...

In [18]:

```
items_per_order_per_user_df.add_to_cart_order.describe()
```

...

In [21]:

```
items_per_order_per_user_df[items_per_order_per_user_df['add_to_cart_order'] >10]
```

...

In [22]:

```
#Understanding the distributions - mean, min, max, median for both the variables to  
#lapply( avg_items_per_order_per_user_df['order_id'].describe(), round, digits=2)  
avg_items_per_order_per_user_df['order_id'].describe()
```

...

In [23]:

```
#number of reordered items per order  
 #(sum for all orders) / total number of orders (avg number of reordered items per o
```

In [24]:

```
#num_reordered_items_per_order_df = ic_df.groupby(['user_id','order_id'],as_index=False)  
reorder_sum_of_all_orders_df = ic_df.groupby(['user_id'],as_index=False)['reordered  
reorder_sum_of_all_orders_df
```

...

In [25]:

```
#total orders per user  
total_orders_per_user_df = ic_df.groupby(['user_id']).order_id.nunique().reset_index  
#total_orders_per_user_df.to_csv("total_order_per_user.csv")  
total_orders_per_user_df
```

...

In [26]:

```
#avg number of reordered items per order per user  
avg_reordered_items_per_order_per_user = np.ceil(reorder_sum_of_all_orders_df['reor  
  
avg_reordered_items_per_order_per_user = pd.DataFrame(avg_reordered_items_per_order_  
avg_reordered_items_per_order_per_user
```

...

In [27]:

```
#determining mean and max of 'days_since_prior_order' for a user  
  
mean_days_since_prior_df = ic_df.groupby(['user_id'],as_index=False)['days_since_pri  
mean_days_since_prior_df.loc[:, 'days_since_prior'] = np.ceil(mean_days_since_prior_  
  
max_days_since_prior_df = ic_df.groupby(['user_id'],as_index=False)['days_since_pri  
max_days_since_prior_df.loc[:, 'days_since_prior'] = np.ceil(max_days_since_prior_d  
  
max_days_since_prior_df
```

...

In [28]:

```
#determining mode of 'order_dow' for a user
```

```
mode_order_dow = ic_df.groupby(['user_id']).agg(lambda x:x.value_counts().index[0])  
mode_order_dow
```

...

In [31]:

```
mode_order_dow1 = mode_order_dow.reset_index()  
mode_order_dow2 = pd.DataFrame(mode_order_dow1['order_dow'])  
  
mode_order_hour_of_day = pd.DataFrame(mode_order_dow1['order_hour_of_day'])  
mode_order_hour_of_day
```

...

In [32]:

```
ic_df.groupby('user_id').first().reset_index()
```

...

In [33]:

```
ic_train_orders_df = ic_df[ic_df['eval_set'] == 'train']  
ic_train_orders_df
```

...

In [34]:

```
#ic_train_orders_df.groupby(['user_id'],as_index=False)['add_to_cart_order'].max()  
  
ic_inter_df = ic_train_orders_df.groupby(['user_id'])[["order_id", "department_id",
```

In [35]:

```
ic_inter_df.shape
```

...

In [36]:

```
ic_inter_df.head()
```

...

In [37]:

```
#concat avg_orders_per_user to the main dataframe ic_inter_df
```

```
ic_inter_df1 = pd.concat([ic_inter_df, total_orders_per_user_df['order_id']], axis=1)
ic_inter_df1
```

...

In [38]:

```
#concat items_per_order_per_user_df to the dataframe - ic_inter_df1
```

```
ic_inter_df2 = pd.concat([ic_inter_df1, items_per_order_per_user_df['add_to_cart_order']], axis=1)
ic_inter_df2
```

...

In [39]:

```
#concat avg_reordered_items_per_order_per_user to the dataframe - ic_inter_df2
```

```
ic_inter_df3 = pd.concat([ic_inter_df2, avg_reordered_items_per_order_per_user], axis=1)
ic_inter_df3
```

...

In [40]:

```
#concat mean_days_since_prior_df to the dataframe - ic_inter_df3
```

```
ic_inter_df4 = pd.concat([ic_inter_df3, mean_days_since_prior_df], axis=1)
ic_inter_df4
```

...

In [41]:

```
#concat max_days_since_prior_df to the dataframe - ic_inter_df4
```

```
ic_inter_df5 = pd.concat([ic_inter_df4, max_days_since_prior_df], axis=1)
ic_inter_df5
```

...

In [42]:

```
#concat mode_order_dow2 to the dataframe - ic_inter_df5
```

```
ic_inter_df6 = pd.concat([ic_inter_df5, mode_order_dow2], axis=1)
ic_inter_df6
```

...

In [43]:

```
#concat mode_order_hour_of_day to the dataframe - ic_inter_df6

ic_df_final = pd.concat([ic_inter_df6, mode_order_hour_of_day], axis=1)
ic_df_final
```

...

In [72]:

```
ic_df_final.head()
```

...

In [64]:

```
ic_df_final.columns.values
```

...

In [66]:

```
#ic_df_final.drop(ic_df_final.columns[[18,20]], axis =1, inplace = True)
ic_df_final.columns.values
```

...

In [67]:

```
ic_df_final.shape
```

...

In [68]:

```
new_columns = ic_df_final.columns.values

new_columns[14] = 'avg_orders'
new_columns[15] = 'avg_items_per_order'
new_columns[16] = 'avg_reordered_items_per_order'
new_columns[17] = 'avg_days_since_prior'
new_columns[18] = 'max_days_since_prior'
new_columns[19] = 'mode_order_dow'
new_columns[20] = 'mode_order_hour_of_day'
```

```
ic_df_final.columns = new_columns
ic_df_final.head()
```

...

In [69]:

```
ic_df_final.columns.values
```

...

In [75]:

```
user_id_df = ic_inter_df6.iloc[:,0]
user_id_df
```

...

In [76]:

```
ic_df_final1 = pd.concat([ic_df_final, user_id_df], axis=1)
ic_df_final1
```

...

In [77]:

```
ic_df_final1.columns.values
```

...

In [78]:

```
#changing the order of the columns
instacart_df = ic_df_final1[['user_id', 'order_id', 'department_id', 'department', 'a
    'product_id', 'product_name', 'eval_set', 'order_number',
    'order_dow', 'order_hour_of_day', 'days_since_prior',
    'add_to_cart_order', 'reordered', 'avg_orders',
    'avg_items_per_order', 'avg_reordered_items_per_order',
    'avg_days_since_prior', 'max_days_since_prior', 'mode_order_dow',
    'mode_order_hour_of_day']]
instacart_df.head()
```

...

In [256]:

```
instacart_df.to_csv("instacart_snapshot.csv")
```


In [79]:

```
#Creating target variable power_user
#(orders per user > 15 and items per order per user > 10) -> count how many labels

def f(row):
    if ((row['avg_orders'] > 10) & (row['avg_items_per_order'] > 9)):
        val = 1
    else:
        val = 0
    return val

instacart_df['power_user'] = instacart_df.apply(f, axis=1)
```

...

In [117]:

```
len(instacart_df[instacart_df['power_user'] ==1])
```

...

In [121]:

```
instacart_df['avg_orders'].hist()
```

...

Functions

In [222]:

```
def plotting_roc(fpr_val,tpr_val,roc_auc_val):
    plt.title('Receiver Operating Characteristic')
    plt.plot(fpr_val, tpr_val, 'b', label = 'AUC = %0.2f' % roc_auc_val)
    plt.legend(loc = 'lower right')
    plt.plot([0, 1], [0, 1], 'r--')
    plt.xlim([0, 1])
    plt.ylim([0, 1])
    plt.ylabel('True Positive Rate')
    plt.xlabel('False Positive Rate')
    plt.show()
    plt.savefig('roc1.png')

def roc_for_thresholds(y,scores):
    # Compute ROC curve and ROC area for each class

    # Use roc_curve to return the TPR and FPR rates at various thresholds
    fpr, tpr, thresholds = roc_curve(y, scores, pos_label=1)

    fpr_df = pd.DataFrame(fpr)
    tpr_df = pd.DataFrame(tpr)
    threshold_df = pd.DataFrame(thresholds)
```

```

tpr_fpr_df = pd.concat([fpr_df,tpr_df],axis =1)
metrics_df = pd.concat([tpr_fpr_df,threshold_df],axis =1)
#print(metrics_df)

print('FPR:' + str(fpr))
print('TPR:' + str(tpr))
print('Thresholds:' + str(thresholds))
# Plot our ROC curve!
plt.plot(fpr, tpr)
plt.xlabel('FPR')
plt.ylabel('TPR')

return(metrics_df)

```

```

def plot_response(k,knn_accuracy):
    #print(len(k))
    plt.plot(k,knn_accuracy,lw=2)

    plt.legend(['knn accuracy'])

    plt.xlabel('k')
    plt.ylabel('accuracy')
    plt.title('Accuracy response to k')
    plt.show()

```

```

def plot_learning_curve(estimator, title, X, y, ylim=None, cv=None,
                        n_jobs=1, train_sizes=np.linspace(.1, 1.0, 5)):
    plt.figure()
    plt.title(title)
    if ylim is not None:
        plt.ylim(*ylim)
    plt.xlabel("Training examples")
    plt.ylabel("Score")
    train_sizes, train_scores, test_scores = learning_curve(
        estimator, X, y, cv=cv, n_jobs=n_jobs, train_sizes=train_sizes)

    train_scores_mean = np.mean(train_scores, axis=1)
    test_scores_mean = np.mean(test_scores, axis=1)
    plt.grid()

    plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
             label="Training score")
    plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
             label="Cross-validation score")

    plt.legend(loc="best")
    return plt

```

```

def plot_confusion_matrix(cm, classes,
                          normalize=False,
                          title='Confusion matrix',
                          cmap=plt.cm.Blues):

```

```
cmap=plt.cm.Blues):
```

```
"""
```

```
This function prints and plots the confusion matrix.  
Normalization can be applied by setting `normalize=True`.
```

```
"""
```

```
plt.imshow(cm, interpolation='nearest', cmap=cmap)  
plt.title(title)  
plt.colorbar()  
tick_marks = np.arange(len(classes))  
plt.xticks(tick_marks, classes, rotation=45)  
plt.yticks(tick_marks, classes)
```

```
if normalize:
```

```
    cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]  
    print("Normalized confusion matrix")
```

```
else:
```

```
    print('Confusion matrix, without normalization')
```

```
print(cm)
```

```
thresh = cm.max() / 2.
```

```
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):  
    plt.text(j, i, cm[i, j],  
            horizontalalignment="center",  
            color="white" if cm[i, j] > thresh else "black")
```

```
plt.tight_layout()  
plt.ylabel('True label')  
plt.xlabel('Predicted label')
```

```
def logistic_regression(ind_var_train,dep_var_train,ind_var_test,dep_var_test):
```

```
    logr = LogisticRegression()  
    logr.fit(ind_var_train,dep_var_train)  
    y_pred = logr.predict(ind_var_test)  
    y_pred_prob = logr.predict_proba(ind_var_test)  
    y_pred_prob = y_pred_prob[:,1]
```

```
    accuracy = accuracy_score(dep_var_test,y_pred)  
    #f1_score = f1_score(dep_var_test, y_pred,average='weighted')  
    recall = recall_score(dep_var_test, y_pred,average='weighted')  
    precision = precision_score(dep_var_test, y_pred,average='weighted')
```

```
    metrics_df = roc_for_thresholds(dep_var_test,y_pred_prob)  
    fpr, tpr, threshold = roc_curve(dep_var_test,y_pred_prob)  
    roc_auc = auc(fpr, tpr)
```

```
# Check trained model intercept  
print("Intercept :",logr.intercept_)
```

```
# Check trained model coefficients  
print("Coefficients :",logr.coef_)
```

```
joblib.dump(logr, 'logistic_model.pkl', protocol=2)
```

```
joblib.dump(logr, instcat_model.pkl, protocol=2)
```

```
plotting_roc(fpr, tpr, roc_auc)
```

```
print("Accuracy :", accuracy)
```

```
#print(f1_score)
```

```
print("Recall :", recall)
```

```
print("Precision :", precision)
```

```
#print(fpr)
```

```
#print(tpr)
```

```
print("ROC_AUC :", roc_auc)
```

```
# View summary of common classification metrics
```

```
print(classification_report(dep_var_test, y_pred))
```

```
# Compute confusion matrix
```

```
class_names = [1, 0]
```

```
cnf_matrix = confusion_matrix(dep_var_test, y_pred)
```

```
np.set_printoptions(precision=2)
```

```
# Plot non-normalized confusion matrix using matplotlib
```

```
plt.figure()
```

```
plot_confusion_matrix(cnf_matrix, classes=class_names,  
                      title='Confusion matrix')
```

```
# Plot normalized confusion matrix
```

```
#plt.figure()
```

```
#plot_confusion_matrix(cnf_matrix, classes=class_names, normalize=True,  
#                      title='Normalized confusion matrix')
```

```
plt.show()
```

```
#Plot confusion matrix using Seaborn
```

```
cm = confusion_matrix(dep_var_test, y_pred)
```

```
df_cm = pd.DataFrame(cm, index = ['True (positive)', 'True (negative)'])
```

```
df_cm.columns = ['Predicted (positive)', 'Predicted (negative)']
```

```
sns.heatmap(df_cm, annot=True, fmt="d")
```

```
return accuracy, recall, precision, metrics_df
```

```
def logistic_regression_cv(ind_var, dep_var, cv):
```

```
    logr_cv = LogisticRegression()
```

```
    accuracy = cross_val_score(logr_cv, ind_var, dep_var, cv=cv, scoring='r2')
```

```
    precision = cross_val_score(logr_cv, ind_var, dep_var, cv=cv, scoring='precision')
```

```
    recall = cross_val_score(logr_cv, ind_var, dep_var, cv=cv, scoring='recall')
```

```
    print("Accuracy :", accuracy.mean())
```

```
    print("Precision :", precision.mean())
```

```
    print("Recall :", recall.mean())
```

```

return accuracy, recall, precision

def logistic_regression_holdout(ind_var_train,dep_var_train,ind_var_test,dep_var_test,cv):
    # Create the hyperparameter grid
    c_space = np.logspace(-5, 8, 15)
    param_grid = {'C': c_space, 'penalty': ['l1', 'l2']}

    # Instantiate the logistic regression classifier: logreg
    logreg = LogisticRegression()

    # Instantiate the GridSearchCV object: logreg_cv
    logreg_cv = GridSearchCV(logreg, param_grid, cv=cv)

    # Fit it to the training data
    logreg_cv.fit(ind_var_train, dep_var_train)

    # Print the optimal parameters and best score
    print("Tuned Logistic Regression Parameter: {}".format(logreg_cv.best_params_))
    print("Tuned Logistic Regression Accuracy: {}".format(logreg_cv.best_score_))

def logistic_regression_poly(ind_var_train,dep_var_train,ind_var_test,dep_var_test,cv):
    #degree = 3
    # Generate the model type with make_pipeline
    # This tells it the first step is to generate 3rd degree polynomial features in
    # a linear regression on the resulting features
    est = make_pipeline(PolynomialFeatures(degree), LogisticRegression())
    # Fit our model to the training data
    est.fit(ind_var_train, dep_var_train)
    #est.score(X_test,y_test)
    y_pred = est.predict(ind_var_test)
    y_pred_prob = est.predict_proba(ind_var_test)
    y_pred_prob = y_pred_prob[:,1]

    accuracy = accuracy_score(dep_var_test,y_pred)
    #f1_score = f1_score(dep_var_test, y_pred,average='weighted')
    recall = recall_score(dep_var_test, y_pred,average='weighted')
    precision = precision_score(dep_var_test, y_pred,average='weighted')

    fpr, tpr, threshold = roc_curve(dep_var_test,y_pred_prob)
    roc_auc = auc(fpr, tpr)

    plotting_roc(fpr,tpr,roc_auc)

    print("Accuracy :", accuracy)
    print("Recall :", recall)
    print("Precision :", precision)
    print("ROC_AUC :", roc_auc)

    return accuracy, recall, precision

def gaussian_nb(ind_var_train,dep_var_train,ind_var_test,dep_var_test):

```

```

gnb = GaussianNB()
gnb.fit(ind_var_train,dep_var_train)
y_pred = gnb.predict(ind_var_test)
y_pred_prob = gnb.predict_proba(ind_var_test)
y_pred_prob = y_pred_prob[:,1]

accuracy = accuracy_score(dep_var_test,y_pred)
#f1_score = f1_score(dep_var_test, y_pred,average='weighted')
recall = recall_score(dep_var_test, y_pred,average='weighted')
precision = precision_score(dep_var_test, y_pred,average='weighted')

fpr, tpr, threshold = roc_curve(dep_var_test,y_pred_prob)
roc_auc = auc(fpr, tpr)

plotting_roc(fpr,tpr,roc_auc)

print("Accuracy :", accuracy)
print("Recall :", recall)
print("Precision :", precision)
print("ROC_AUC :", roc_auc)

return accuracy, recall, precision

```

```

def support_vector_machine(ind_var_train,dep_var_train,ind_var_test,dep_var_test):
    model_svm = svm.SVC(kernel='rbf',probability=True)
    model_svm.fit(ind_var_train,dep_var_train)
    y_pred = model_svm.predict(ind_var_test)
    y_pred_prob = model_svm.predict_proba(ind_var_test)
    y_pred_prob = y_pred_prob[:,1]

    accuracy = accuracy_score(dep_var_test,y_pred)
    #f1_score = f1_score(dep_var_test, y_pred,average='weighted')
    recall = recall_score(dep_var_test, y_pred,average='weighted')
    precision = precision_score(dep_var_test, y_pred,average='weighted')

    fpr, tpr, threshold = roc_curve(dep_var_test,y_pred_prob)
    roc_auc = auc(fpr, tpr)

    plotting_roc(fpr,tpr,roc_auc)

    print("Accuracy :", accuracy)
    print("Recall :", recall)
    print("Precision :", precision)
    print("ROC_AUC :", roc_auc)

    return accuracy, recall, precision

```

```

def decision_tree(ind_var_train,dep_var_train,ind_var_test,dep_var_test):
    dt = DecisionTreeClassifier()
    dt.fit(ind_var_train,dep_var_train)
    y_pred = dt.predict(ind_var_test)

```

```

y_pred_prob = dt.predict_proba(ind_var_test)
y_pred_prob = y_pred_prob[:,1]

accuracy = accuracy_score(dep_var_test,y_pred)
#f1_score = f1_score(dep_var_test, y_pred,average='weighted')
recall = recall_score(dep_var_test, y_pred,average='weighted')
precision = precision_score(dep_var_test, y_pred,average='weighted')

fpr, tpr, threshold = roc_curve(dep_var_test,y_pred_prob)
roc_auc = auc(fpr, tpr)

plotting_roc(fpr,tpr,roc_auc)

print("Accuracy :", accuracy)
print("Recall :", recall)
print("Precision :", precision)
print("ROC_AUC :", roc_auc)

return accuracy, recall, precision

```

```

def random_forest(ind_var_train,dep_var_train,ind_var_test,dep_var_test):
    rf = RandomForestClassifier()
    rf.fit(ind_var_train,dep_var_train)
    y_pred = rf.predict(ind_var_test)

    y_pred_prob = rf.predict_proba(ind_var_test)
    y_pred_prob = y_pred_prob[:,1]

    accuracy = accuracy_score(dep_var_test,y_pred)
    #f1_score = f1_score(dep_var_test, y_pred,average='weighted')
    recall = recall_score(dep_var_test, y_pred,average='weighted')
    precision = precision_score(dep_var_test, y_pred,average='weighted')

    fpr, tpr, threshold = roc_curve(dep_var_test,y_pred_prob)
    roc_auc = auc(fpr, tpr)

    plotting_roc(fpr,tpr,roc_auc)

    print("Accuracy :", accuracy)
    print("Recall :", recall)
    print("Precision :", precision)
    print("ROC_AUC :", roc_auc)

    return accuracy, recall, precision

```

```

def gradient_boost(ind_var_train,dep_var_train,ind_var_test,dep_var_test):
    gb = GradientBoostingRegressor()
    gb.fit(ind_var_train,dep_var_train)
    y_pred = gb.predict(ind_var_test)

    y_pred_prob = gb.predict_proba(ind_var_test)
    y_pred_prob = y_pred_prob[:,1]

    accuracy = accuracy_score(dep_var_test,y_pred)
    #f1_score = f1_score(dep_var_test, y_pred,average='weighted')
    recall = recall_score(dep_var_test, y_pred,average='weighted')
    precision = precision_score(dep_var_test, y_pred,average='weighted')

    fpr, tpr, threshold = roc_curve(dep_var_test,y_pred_prob)
    roc_auc = auc(fpr, tpr)

    plotting_roc(fpr,tpr,roc_auc)

    print("Accuracy :", accuracy)
    print("Recall :", recall)
    print("Precision :", precision)
    print("ROC_AUC :", roc_auc)

    return accuracy, recall, precision

```

```

#f1_score = f1_score(dep_var_test, y_pred,average='weighted')
recall = recall_score(dep_var_test, y_pred,average='weighted')
precision = precision_score(dep_var_test, y_pred,average='weighted')

fpr, tpr, threshold = roc_curve(dep_var_test,y_pred_prob)
roc_auc = auc(fpr, tpr)

plotting_roc(fpr,tpr,roc_auc)

print("Accuracy :", accuracy)
print("Recall :", recall)
print("Precision :", precision)
print("ROC_AUC :", roc_auc)

```

In [30]:

```
final_cols = ['avg_reordered_items_per_order', 'avg_days_since_prior', 'mode_order_c
```

Creating a holdout set for finally checking the model performance

In []:

```

# X = X.loc[:, final_cols]
#X_holdout = X_holdout.loc[:, final_cols]
# final_model = Lasso(alpha = final_alpha)
# final_fit = final_model.fit(X, y)

```

Determining feature importance

1. Checking correlations

In [83]:

```
instacart_df.corr()
```


...

In [213]:

```
ic_corr = instacart_df.iloc[:,15:].corr()  
sns.heatmap(ic_corr,  
            xticklabels=ic_corr.columns.values,  
            yticklabels=ic_corr.columns.values)
```

Out[213]:

<matplotlib.axes._subplots.AxesSubplot at 0x12b6ba160>




In [214]:

```
sns.pairplot(ic_corr)
```

Out[214]:

<seaborn.axisgrid.PairGrid at 0x12afb2e48>



2. Using Random Forest

In []:

In [84]:

```
X = instacart_df.ix[:,['avg_reordered_items_per_order', 'avg_days_since_prior', 'mode_order_dow', 'mode_order']]

#X = ic_inter_df5.iloc[:,9:20]
y = instacart_df.iloc[:, -1]

X.head()
```

Out[84]:

	avg_reordered_items_per_order	avg_days_since_prior	mode_order_dow	mode_order
0	5.0	20.0	4	8
1	7.0	19.0	1	11
2	4.0	13.0	0	18
3	7.0	14.0	0	18
4	5.0	23.0	1	0

In [85]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
rf = RandomForestRegressor(n_estimators=200)
rf.fit(X_train, y_train)
rf.score(X_test, y_test)

imp = rf.feature_importances_

imp = pd.DataFrame(np.array(imp).T, columns = ['imp'], index = X.columns)
imp.sort_values('imp', ascending = False, inplace = True)
#imp.to_csv("important_features.csv")
print(imp)
```

	imp
avg_reordered_items_per_order	0.518378
avg_days_since_prior	0.206350
mode_order_hour_of_day	0.170200
mode_order_dow	0.105072

3. Using Multinomial Naive Bayes

In [86]:

```
# Create the model
mnb = MultinomialNB()

# Fit the model to the training data
mnb.fit(X_train, y_train)
# Score the model against the test data
mnb.score(X_test, y_test)

#mnb.feature_log_prob_
```

Out[86]:

0.83161447590474402

Modeling

Logistic Regression with variables 'avg_reordered_items_per_order', 'avg_days_since_prior', 'mode_order_dow', 'mode_order_hour_of_day'

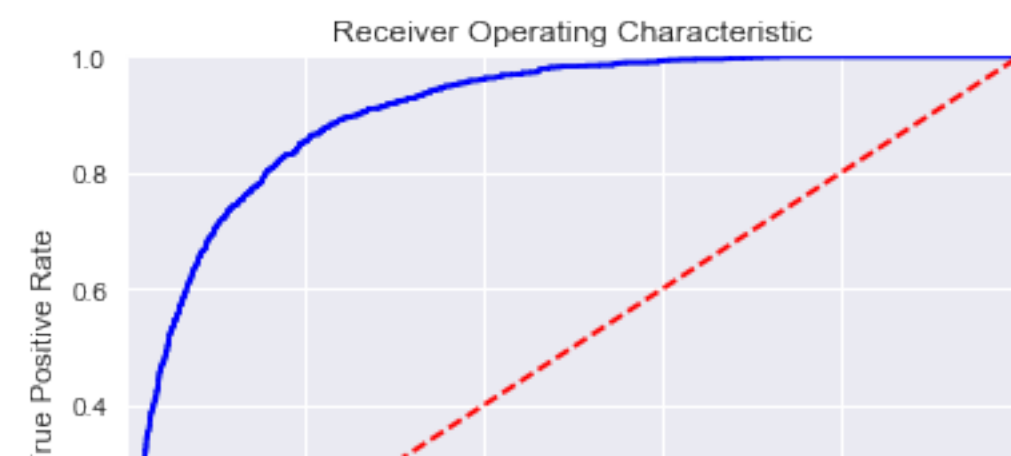
In [223]:


```
X = instacart_df.ix[:,['avg_reordered_items_per_order', 'avg_days_since_prior', 'mode_order_dow', 'mode_order_hour_of_day']]
y = instacart_df.iloc[:, -1]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

accuracy, recall, precision, metrics_df = logistic_regression(X_train, y_train, X_test, y_test)
```

```
FPR:[ 0.00e+00  0.00e+00  2.51e-04 ...,  9.96e-01  9.96e-01  1.0
0e+00]
TPR:[ 7.41e-04  8.15e-03  8.15e-03 ...,  1.00e+00  1.00e+00  1.0
0e+00]
Thresholds:[ 1.    1.    1.    ...,  0.01  0.01  0.   ]
Intercept : [-2.83]
Coefficients : [[ 0.52 -0.09 -0.02  0.   ]]
```







Accuracy : 0.848303018939
 Recall : 0.848303018939
 Precision : 0.842033799225
 ROC_AUC : 0.905314270425


	precision	recall	f1-score	support
0	0.87	0.93	0.90	3984
1	0.75	0.60	0.67	1349
avg / total	0.84	0.85	0.84	5333

Confusion matrix, without normalization

```
[[3718  266]
 [ 543  806]]
```

<matplotlib.figure.Figure at 0x12da0c860>





The model produced a negative intercept value and a weight of 0.52 on avg_reordered_items_per_order, -0.09 on avg_days_since_prior, -0.02 on mode_order_dow and 0. on mode_order_hour_of_day

Logistic Regression with variables 'avg_reordered_items_per_order', 'avg_days_since_prior', 'mode_order_dow', 'mode_order_hour_of_day' and k fold cross validation

In [88]:

```
accuracy, recall, precision = logistic_regression_cv(X, y, 15)
```

Accuracy : 0.2159576962
Precision : 0.771091481707
Recall : 0.592656219264


Learning curve for logistic regression

In [109]:

```
title = "Learning Curves (Logistic Regression)"
# Cross validation with 100 iterations to get smoother mean test and train
# score curves, each time with 20% data randomly selected as a validation set.
cv = ShuffleSplit(n_splits=100, test_size=0.2, random_state=4444)

estimator = LogisticRegression()
plot_learning_curve(estimator, title, X, y, ylim=(0.7, 1.01), cv=cv, n_jobs=4)

plt.show()
```



K Nearest Neighbors

In [97]:

```
#Try it with a lot of different k values (number of neighbors), from 1 to 20,  
#and on the test set calculate the accuracy (number of correct predictions / number
```

```
accuracy = []  
accuracy_index = []  
  
for k in range(1,21):  
    knn = KNeighborsClassifier(n_neighbors=k)  
  
    knn.fit(X_train,y_train)  
    y_pred = knn.predict(X_test)  
    accuracy.append(accuracy_score(y_test,y_pred))  
    accuracy_index.append(k)  
  
print(accuracy)  
print(accuracy_index)  
max_accuracy = max(accuracy)  
print(max_accuracy)  
max_acc_k = accuracy.index(max_accuracy)+1  
print(max_acc_k)
```

```
[0.8151134445902869, 0.82542658916182265, 0.8385524095255954, 0.843240  
20251265708, 0.84830301893868365, 0.85055315957247324, 0.8529908119257  
4533, 0.85467841740108752, 0.85449090568160513, 0.85767860491280701, 0  
.85617851115694732, 0.8589911869491843, 0.85767860491280701, 0.8599287  
4554659671, 0.85617851115694732, 0.8606787924245265, 0.859178698668666  
81, 0.8588036752297018, 0.8616163510219389, 0.860866304144009]  
[1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20  
]  
0.861616351022  
19
```

In [284]:

```
k=19
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1, random_state=42)
knn_19 = KNeighborsClassifier(n_neighbors=k)
knn_19.fit(X_train,y_train)
y_pred = knn_19.predict(X_test)

# Compute confusion matrix
class_names =[1,0]
cnf_matrix = confusion_matrix(y_test, y_pred)
np.set_printoptions(precision=2)

# Plot non-normalized confusion matrix using matplotlib
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=[1,0],
                      title='Confusion matrix')

# Plot normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=class_names, normalize=True,
                      title='Normalized confusion matrix')
plt.savefig('knn_conf_mat.png')
plt.show()
plt.savefig('knn_conf_mat.jpg')

accuracy = accuracy_score(y_test, y_pred)
#f1_score = f1_score(dep_var_test, y_pred,average='weighted')
recall = recall_score(y_test, y_pred,average='weighted')
precision = precision_score(y_test, y_pred,average='weighted')

fpr, tpr, threshold = roc_curve(y_test, y_pred)
roc_auc = auc(fpr, tpr)


print("Accuracy :", accuracy)
print("Recall :", recall)
print("Precision :", precision)
print("ROC_AUC :", roc_auc)
```

...

KNN - Accuracy as a function of k

In [98]:

```
plot_response(accuracy_index, accuracy)
```




Learning Curve for KNN where k is the one with maximum accuracy

In [111]:

```
title = "Learning Curves (KNN)"
# Cross validation with 100 iterations to get smoother mean test and train
# score curves, each time with 20% data randomly selected as a validation set.
cv = ShuffleSplit(n_splits=100, test_size=0.2, random_state=1000)

estimator = KNeighborsClassifier(n_neighbors=19)
plot_learning_curve(estimator, title, X, y, ylim=(0.7, 1.01), cv=cv, n_jobs=4)

plt.show()
```



Gaussian Naive Bayes

In [90]:

```
accuracy, recall, precision = gaussian_nb(X_train,y_train,X_test,y_test)
```




Accuracy : 0.854490905682
Recall : 0.854490905682
Precision : 0.849164185812
ROC_AUC : 0.897884439909

Support Vector Machine

In [91]:

```
accuracy, recall, precision = support_vector_machine(X_train,y_train,X_test,y_test)
```



Accuracy : 0.857678604913
Recall : 0.857678604913
Precision : 0.854305531346
ROC_AUC : 0.858608916764

Decision Tree

In [92]:

```
accuracy, recall, precision = decision_tree(X_train,y_train,X_test,y_test)
```



Accuracy : 0.818301143821
Recall : 0.818301143821
Precision : 0.813761851879
ROC_AUC : 0.768207559668

Random Forest

In [93]:

```
accuracy, recall, precision = random_forest(X_train,y_train,X_test,y_test)
```



Accuracy : 0.840052503281

Recall : 0.840052503281


Precision : 0.83649864831

ROC_AUC : 0.869188391818


In [217]:

```
accuracy, recall, precision = logistic_regression_poly(X_train,y_train,X_test,y_test)
plot_learning_curve(est, title, X, y, ylim=(0.7, 1.01), cv=cv, n_jobs=4)

plt.show()
```



Accuracy : 0.858241140071
Recall : 0.858241140071
Precision : 0.855683687216
ROC_AUC : 0.901685224962




In [205]:

```
degree = 3
# Generate the model type with make_pipeline
# This tells it the first step is to generate 3rd degree polynomial features in the
# a linear regression on the resulting features
est = make_pipeline(PolynomialFeatures(degree), LogisticRegression())
# Fit our model to the training data
est.fit(X_train, y_train)
est.score(X_test, y_test)

plot_learning_curve(est, title, X, y, ylim=(0.7, 1.01), cv=cv, n_jobs=4)

plt.show()
```



Hold-out set in practice: Classification

Evaluating a model with tuned hyperparameters on a hold-out set.

In addition to C, logistic regression has a 'penalty' hyperparameter which specifies whether to use 'l1' or 'l2' regularization. Your job in this exercise is to create a hold-out set, tune the 'C' and 'penalty' hyperparameters of a logistic regression classifier using GridSearchCV on the training set, and then evaluate its performance against the hold-out set.

In [94]:

```
cv = 5
logistic_regression_holdout(X_train, y_train, X_test, y_test, cv)
```

...

Comparison of models using cross validation

In []:

```
import sys
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.dummy import DummyClassifier

bestKValue = 19
models = {}
models = {'logres': LogisticRegression(), # Takes about 1 second elapsed time
          'knn with K=%d' % bestKValue : KNeighborsClassifier(n_neighbors=bestKValue),
          'gaussianNB': GaussianNB(),
          'random forest': RandomForestClassifier(),
          'decision tree': DecisionTreeClassifier(),
          'svm': svm.SVC(kernel='rbf',probability=True),
          'baseline' : DummyClassifier(strategy='stratified')}
#          'randomforest with nrEst=%d and maxFeat=%d' % (bestNrEst,bestMaxFeat): RandomForestClassifier(n_estimators=bestNrEst,max_features=bestMaxFeat)

scorerType = 'roc_auc'

nrCrossValidationFolds=10
# We MUST shuffle, because the data seem to be somehow ordered
cvGenerator = KFold(len(X), n_folds=nrCrossValidationFolds, shuffle=True)

fig = plt.figure(1,(9,6))
plt.title("%d-fold cross-validation %s scores for various model types" % (nrCrossValidationFolds,scorerType))
plt.xlabel("Fold #")
plt.ylabel(scorerType)
plt.grid()

for modelName, model in models.items():
#     print >> sys.stderr, "Building %s model ..." % modelName,
#     print >> sys.stderr, "applying it ...",
    scores = cross_val_score(model, X, y=y, scoring=scorerType, cv=cvGenerator, n_jobs=-1)
#     print >> sys.stderr, "done"
    plt.plot(range(1,nrCrossValidationFolds+1), scores, 'o-', label="%s (%2.2f%% +/- %2.2f%%)" % (modelName, scores.mean()*100, scores.std()*100))

plt.legend(loc='best',fontsize = 'large')
plt.savefig('model_comp_kfolds.png')
plt.show()
```

...

ROC curves - Balancing true positives and false positives

In [3]:

```
import mpld3
mpld3.enable_notebook()
```



```

import re
from math import log
import json
import copy

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1, random_state=42)

finalPredictions = {}

for modelName, model in models.items():
    fittedModel = model.fit(X_train, y_train)

    if hasattr(model, "predict_proba"):
        #print >> sys.stderr, "Predicting probabilities for %s model ..." % modelName
        finalPredictions[modelName] = {'train': model.predict_proba(X_train),
                                       'test': model.predict_proba(X_test)}
        #print >> sys.stderr, "done"

def plotROCCurve(figSize=7):
    # Define some CSS to control our custom labels
    css = """
table {
    border-collapse: collapse;
}
th {
    color: #ffffff;
    background-color: #000000;
}
td {
    padding: 2px;
    background-color: #cccccc;
}
table, th, td {
    font-family: Arial, Helvetica, sans-serif;
    border: 1px solid black;
    text-align: right;
}
"""

    jsonROCDData = {}
    rocCurveFigure, ax = plt.subplots(figsize=(figSize, figSize))
    ax.grid(True, alpha=0.3)
    for modelName, probs in finalPredictions.items():
        modelNameShort = re.split("\s+", modelName)[0]
        y_probs = [x[1] for x in finalPredictions[modelName]['test']]
        fpr, tpr, thresholds = roc_curve(y_true=y_test,
                                         y_score=y_probs, pos_label=1)

        roc_auc = roc_auc_score(y_true=y_test, y_score=y_probs)
        jsonROCDData[modelNameShort] = {}
        jsonROCDData[modelNameShort]['fpr'] = [x for x in fpr]
        jsonROCDData[modelNameShort]['tpr'] = [x for x in tpr]
        jsonROCDData[modelNameShort]['thresholds'] = [np.asscalar(np.float32(x)) for

```

```

        jsonROCDData[modelNameShort]['roc_auc'] = roc_auc
        points = plt.plot(fpr, tpr, 'x-', label="%s (AUC = %1.2f%%)" % (modelName,
        labels = ["<table><th colspan='2'>%s</th><tr><td>FPR</td><td>%0.1f%%</td></tr>"]
        mpld3.plugins.connect(rocCurveFigure, mpld3.plugins.PointHTMLTooltip(points))
        ax.set_title("ROC curve for prediction of prime user", y=1.06, fontsize=14 + log(figSize))
        ax.set_xlabel("False Positive Rate (FP/FP+TN)", labelpad=15 + log(figSize), fontweight='bold')
        ax.set_ylabel("True Positive Rate (TP/TP+FN)", labelpad=15 + log(figSize), fontweight='bold')
        plt.legend(loc="best", fontsize = 'xx-large')
        plt.show()
        mpld3.save_html(rocCurveFigure, 'ROC_comp')
        #plt.savefig('ROC_comp')
        rocCurveFigure.savefig('ROC_comp2')
        return jsonROCDData

```

```

# Export data to JSON file for visualization in D3.js or similar
jsonROCDData = plotROCCurve(figSize=9)
#print(jsonROCDData)
#with open('d3/ROCCurve.json', 'w') as outfile:
with open('ROCCurve.json', 'w') as outfile:
    json.dump(jsonROCDData, outfile)

#from IPython.html.widgets import interact, fixed
from ipywidgets import interact, interactive, fixed

interact(plotROCCurve, figSize=(5,10))

```

...